

An Introduction to Reinforcement Learning

Snake Your Way Out of a Paper Bag

Frances Buontempo


2025

Talk outline


- Overview of reinforcement learning
 - It's a huge topic so just some parts
- Start simple
 - Move in 1D
- A bit more complicated
 - Move in 2D
- Fun and games
 - Snake!!!!!!!
 - And other arcade games
- (There will be some C++)




- I edit ACCU's Overload magazine
 - <https://accu.org/journals/nonmembers/overload> cover members/
- Programmer (C++ plus some Python and C#) (mostly in finance)
- Author



Introducing C++



Learn C++ by Example



Genetic Algorithms and
Machine Learning for
Programmers

Where am I?

- <https://mastodon.social/@fbuontempo>
- <https://bsky.app/profile/fbuontempo.bsky.social>
- <https://x.com/fbuontempo>
 - (formerly <https://twitter.com/fbuontempo>)
- <https://www.linkedin.com/in/francesbuontempo/>
- <https://buontempoconsulting.blogspot.com/>
 - (Sometimes)

Reinforcement learning (RL)


- A type of machine learning, so part of AI.
- Involves agent(s) acting in an environment
 - Static or dynamic
 - Continuous or discrete
 - Action + reward -> learn
- Nothing to do with ChatGPT
 - But Reinforcement Learning from Human Feedback (**RLHF**)
 - Uses human feedback to fine tune LLMs (supervised fine-tuning).

A simple game:

```
while (!game_over)
    action = pick_action(state)
    reward = act(action, state)
```


Single random agent, static environment

- An agent, shown by a green square, can **act**
 - By moving left or right, but not off the edge of the world
- Reward (score) is -1 for off the left, +1 for out off the right.
- No learning this time
 - Movement stochastic or random



```


void rnd()
{
    constexpr size_t line_length = 5;
    int position = 2;
    std::mt19937 gen(std::random_device{}());
    std::uniform_int_distribution dist(-1, 1);

    draw(world(line_length, position));
    while (position > -1 && position < line_length)
    {
        auto action = dist(gen);
        while (action == 0) { //want -1 or +1 only
            action = dist(gen);
        }
        position += action;
    }
    draw(world(line_length, position)); //world is just a string to display for now
}

```




```
void draw(const std::string & s)
{
    using namespace std::chrono;
    std::cout << "\x1B[2J\x1B[H";
    std::cout << s;
    std::this_thread::sleep_for(1000ms);
}
```



A random walk

- Expect Left/Right : 50/50
- Example run: Lost 44, Won 56
- Average steps: 9.0
- Let's **learn** now

Single agent with model, static environment

- Previously, move was -1 or +1, at random.
- Now we use
 - Pick action (left/right) based on **model**
 - Which can be a fixed action, like go right
 - Or do something random, with probability ϵ (epsilon)
- Hard code action first...
 - Then we'll let the agent find this out.
- Single agent, but has several goes (episodes)

General idea:

```
for_each(episode)
    env.reset()
    while(!game_over)
        action = pick_action(state)
        reward = env.step(action, state)
        // ignored,
        // but could be used to update model
```

```

void fixed_model() {
    constexpr size_t line_length = 5;
    int position = 2;
    std::mt19937 gen(std::random_device{}());
    std::uniform_int_distribution<> dist(-1, 1);
    std::uniform_real_distribution<> prob(0.0, 1.0);
    constexpr double epsilon = 0.1;

```

probability $\varepsilon = 0.1 = \frac{1}{10}$

```


    draw(world(line_length, position));
    while (position > -1 && position < line_length) {
        auto action = prob(gen) < epsilon ?
            dist(gen) : best_fixed_action(position);

        while (action == 0)
            action = dist(gen);
        position += action;
        draw(world(line_length, position));
    }
}


```

$$action = \begin{cases} rnd, & p < \frac{1}{10} \\ best, & x \geq \frac{1}{10} \end{cases}$$

A model (or fixed action)



```
int best_fixed_action(int position)
{
    return 1; // Strong types would be better
}
```

Stats: (almost) everyone's a winner

- Example run: Lost 0, Won 100
- Average steps: 3.28
- The agent only has a 50/50 chance to go left 1/10 times.
- It didn't learn anything though
- We told it a model
- Can it figure this out on its own?

Now let's try to learn

- Previously we used
 - Pick action based on a model
 - Or did something random, with probability ϵ (epsilon)
- Now, we'll pick action for current state and store reward
 - Reward -1 for out on left, +1 for out on right, 0 otherwise
 - `lookup[(state, action)] -> reward`
 - Key: position and action pair
 - Value: action
- Note: Picking an action based on previous rewards
 - Assumes the rewards are stationary
- `std::map<std::pair<int, int>, int> quality;`

General idea:

```
Lookup quality{}  
for_each(episode)  
    env.reset()  
    while(!game_over)  
        action = pick_action(state, quality)  
        reward = env.step(action, state)  
        learn(action, state, reward)
```

```

auto action = prob(gen) < epsilon ?
    dist(gen) : best_action(position, quality);


int best_action(int position,
    std::map<std::pair<int, int>, int> & quality)
{
    auto left = quality[{position, -1}]; //quality not const
    auto right = quality[{position, 1}]; //'cos a it's map

    if (left == right)
        return 0; // to say do a rnd, enum maybe clearer
    if (left > right)
        return -1; // left
    return 1; //right
}

```

```
int reward(int position, size_t line_length)
{
    if (position < 0)
        return -1;
    else if (std::cmp_less(position, line_length))
        return 0;
    return +1;
}
```


```
void learn(int state, int action, int reward
          std::map<std::pair<int, int>, int>& quality)
{
    quality[{state, action}] += reward;
}
```

Stats

- Example run: Lost 11, Won 89
- Average steps: 11.88
- Recall, with the fixed model:
 - Example run: Lost 0, Won 100
 - Average steps: 3.28
- Better than pure random:
 - Example run: Lost 44, Won 56
 - Average steps: 9.0

Quality



```
best_action:
if (quality(left) == quality(right))
    return 0; // ? shrug
if (quality(left) > quality(right))
    return -1; // left
return 1; // right
```


Position	Action	Quality
0	Left	-11
0	Right	0
1	Left	0
1	Right	0
2	Left	0
2	Right	0
3	Left	0
3	Right	0
4	Left	0
4	Right	0
5	Left	0
5	Right	89

So much nothing

- Quality: -11, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 89
- Look at all the 0s!
- It explores a bit
 - Epsilon means it sometimes fails
 - Greedy (best always) instead means it only steps off the left edge once
- Agent has learnt to avoid going left
 - BUT only knows one step right is better at the right hand edge

Quality Learning

- Q-learning
 - Chris Watkins, 1989 PhD
- One type of RL
- Doesn't need a model
- The quality table uses last and next state
 - a % of the next known rewards add to the state
 - meaning newer rewards can be more important
 - so agents can learn in a dynamic env



Crowd your way out of a paper bag

Door Field

5	5	5	5	5	5
4	4	4	4	4	4
3	3	3	3	3	3
2	2	2	2	2	2
1	1	1	1	1	1
0	0	0	0	0	0

<https://www.youtube.com/watch?v=wlsbg5q0hO0&t=5s>

But now the agent learns

- Recap: so much nothing
- Want to **learn**, without of floor field or model
 - The agent will build up this “field” through experience
 - This “field” is called a **policy** in RL: what action to take in a given state
- Want agent to discover to head in a direction
 - So not have lots of 0s in the middle
- One type of RL: Temporal difference (TD)
 - Notes reward in a given state and remembers what it can do in the new state
 - Including potential maximum reward
 - `learn(state, action, reward) -> learn(last_state, action, reward, next_state)`

Q-Learning

- Q-learning is “Off-policy TD Control”
 - Off-policy: estimates the return (total discounted future reward) for state-action pairs assuming a greedy policy were followed
 - SARSA is on-policy (estimates the return for state-action pairs assuming the current policy continues to be followed.)
- Uses the Bellman equation:
- $Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$
- Quality (for last state) becomes
 - previous quality + α (learning rate) times
 - Reward this time
 - + γ (discount/forgetting) times best possible from this state next state
 - Subtract previous quality
- Percolates next state back
 - So the zeros in the middle go

Q-Learning: It's all Greek to me

- **Epsilon** for random action versus best action based on Q-table.
 - Explore/exploit
- **Alpha** for learning
 - Weights the reward plus potential next reward
- **Gamma**
 - AKA “discount factor”, beta (or lambda), or “impatience”
 - If 0, the agent is **myopic**, since it only uses the latest reward
- The Bellman (optimality) equation
 - Dates to 1950s, introduced in “Dynamic Programming” book
 - calls gamma the discount factor
 - https://en.wikipedia.org/wiki/Bellman_equation

General idea:

```
Lookup quality{}  
for_each(episode)  
    env.reset()  
    while(!game_over)  
        action = pick_action(state, quality)  
        reward = env.step(action, state)  
        q_learn(last_state,  
                action, reward, next_state)
```

```

void q_learn(int state, int action, int reward, int next_state,
             std::map<std::pair<int, int>, double>& quality)
{
    constexpr double alpha = 0.1;
    constexpr double gamma = 0.99;
    double predict = quality[{state, action}];
    auto v = quality
        | std::views::filter([next_state](auto kvp) {
            return std::get<0>(kvp.first) == next_state;
        });
    auto potential_best = v.empty() ? 0.0 :
        std::ranges::max_element(v, [](const auto& lhs, const auto& rhs) {
            return lhs.second < rhs.second;
        })->second;
    double target = reward + gamma * potential_best;
    quality[{state, action}] += alpha * (target - predict);
}

```


Stats:

- Mean steps 3.94
- Q-table
- (0, -1), -0.1
- (0, 1), 0.0265353
- (1, -1), 0
- (1, 1), 0.338015
- (2, -1), 0.0553337
- (2, 1), 0.97863
- (3, -1), 0.379363
- (3, 1), 0.989742
- (4, -1), 0.336458
- (4, 1), 0.99997

Recap

- Now mean steps is 3.94
- Average steps: 11.88 for lots of 0s in the middle
- Recall, with the fixed model:
 - Example run: Lost 0, Won 100
 - Average steps: 3.28
- Better than pure random:
 - Example run: Lost 44, Won 56
 - Average steps: 9.0

Demo



Time to go 2D

- Time to go 2D
- And tidy the code
 - Position was an `int`
- Lookup:
 - Was `std::map<std::pair<int, int>, double> quality;`
 - `std::map<std::pair<std::pair<int, int>, int>, double> (!)`
 - `Lookup<pair<Position, Action>, reward>`
- Reward function needs to change too



General idea (unchanged):

```
Lookup quality{}  
for_each(episode)  
    env.reset()  
    while(!game_over)  
        action = pick_action(state, quality)  
        reward = env.step(action, state)  
        q_learn(last_state,  
                action, reward, next_state)
```

```
struct Position {  
    int x;  
    int y;  
    auto operator<=>(const Position&) const = default;  
};
```

```
enum class Action {  
    Shrug,  
    Left,  
    Right,  
    Up,  
    Down  
};
```

```
using Lookup = std::map<std::pair<Position, Action>,  
                        double>;
```

```

double best_possible_reward_given_state(CharState state, LookupCharOnly& quality) {
    auto v = quality | std::views::filter([state](auto kvp) {
        return std::get<0>(kvp.first) == state; });
    return v.empty() ? 0.0 : std::ranges::max_element(v,
        [](const auto& lhs, const auto& rhs) { return lhs.second < rhs.second; }
    )->second;
}

void q_learn(CharState state, Action action, double reward,
    CharState next_state, LookupCharOnly& quality) {
    constexpr double alpha = 0.1;
    constexpr double gamma = 0.99;
    const double predict = quality[{state, action}];
    const double potential_best = best_possible_reward_given_state(next_state, quality);
    double target = reward + gamma * potential_best;
    quality[{state, action}] += alpha * (target - predict);
}

```

```

Action pick_best_action(CharState state, LookupCharOnly& quality) {
    static std::mt19937 gen(std::random_device{}());
    std::vector<double> weights;
    const std::vector<Action> actions{Action::Up, Action::Down, Action::Left, Action::Right};

    for (auto act : actions)
        weights.push_back(quality[{ state, act }]);

    auto [smallest_pos, biggest_pos] = std::minmax_element(weights.cbegin(), weights.cend());
    auto smallest = *smallest_pos;
    auto biggest = *biggest_pos;

    std::transform(weights.begin(), weights.end(), weights.begin(),
        [biggest](const double& x) { return x == biggest ? 1.0 : 0.0; }
    );
    std::discrete_distribution<int> dd{ weights.cbegin(), weights.cend() };
    return actions[dd(gen)];
}


```

Or easier to see:

- Pick best action
 - Or random choice if two or more equally good
- Don't forget, this only happens sometimes, based on epsilon:

```
auto action = prob(gen) < epsilon ?  
    static_cast<Action>(dist(gen)) :  
    pick_best_action(position, quality);
```

Out of a paper bag!



Stats

Won 20, lost 30

Q-table

(0, 0: Left), -0.1

(0, 0: Down), -0.1

(0, 1: Left), -0.1

(0, 2: Left), -0.19

(0, 3: Left), -0.271

(0, 4: Up), 0.499001

(1, 0: Down), -0.3439

(1, 4: Right), 0.0296703

(1, 4: Up), 0.3994

(2, 0: Down), -0.271

(2, 4: Up), 0.697903

(3, 0: Down), -0.19

(4, 0: Right), -0.3439

(4, 0: Down), -0.19

(4, 1: Right), -0.271

(4, 2: Right), -0.1


(4, 3: Right), -0.19

(4, 4: Right), -0.0703297

(4, 4: Up), 0.3994

What about an obstacle?

- We want to do snake
 - eventually
- Let's start with something to avoid, like 'O' somewhere
 - Then turn it into an apple to eat
 - So the single + becomes a snake
 - (But we'll have to deal with dynamic env then)




Some obstacles

```
std::deque<Position> obstacles{}; // Demo {1,1}

bool game_over() const {
    auto it = std::ranges::find(obstacles, pos);
    return it != obstacles.end()
        || pos.x < 0 || pos.x >= max_x || pos.y < 0 || pos.y >= max_y;
}

int step(Action action) {
    pos = perform_action(pos, action);
    if (game_over())
    {
        return pos.y >= max_y ? 1 : -1; // Allowed out of bag - reward in that case
    }
    return 0;
}
```



One obstacle, greedy action

Won 16, lost 34

(0, 0: Left), -0.1

(0, 0: Down), -0.1

(0, 1: Right), -0.1

(0, 2: Left), -0.271

(0, 3: Left), -0.1

(0, 4: Left), -0.1

(1, 0: Up), -0.1

(1, 0: Down), -0.271

(1, 2: Down), -0.468559

(1, 4: Up), 0.3994

(2, 0: Down), -0.271

(2, 1: Left), -0.468559

(2, 4: Up), 0.697903

(3, 0: Down), -0.1

(4, 0: Right), -0.1

(4, 0: Down), -0.19

(4, 1: Right), -0.1

(4, 2: Right), -0.1

(4, 3: Right), -0.1

(4, 4: Up), 0.499001

Left: -0.1	Up: 0.3994	Up: 0.697903		Up: 0.499001
Left: -0.1				Right: -0.1
Left: -0.271	Down: -0.268559			Right: -0.1
Right: -0.1		Left: -0.468559		Right: -0.1
Left: -0.1, Down -0.1	Up: -0.1, Down -0.271	Down: -0.271	Down: -0.1	Right: -0.1, Down: -0.19

Dynamic env

- What if the environment is dynamic?
 - Change the obstacle to an apple
 - New an apple appears
 - Agent needs to forget old position
- Snake will grow when it eats the apple
 - But we'll just respawn the apple first
 - And grow the snake later
 - New reward function
 - And new danger
- We'll try the same lookup for the Q-table
 - Then a new lookup state



General idea (unchanged, apart from env):

```
Lookup quality{}  
for_each(episode)  
    env.reset()  
    while(!game_over)  
        action = pick_action(state, quality)  
        reward = env.step(action, state)  
        q_learn(last_state,  
                action, reward, next_state)
```

Env changes: no proper snake (yet)

```
int step(Action action) {  
    pos = perform_action(pos, action);  
    if (game_over()) {  
        return pos.y >= max_y ? 1 : -1;  
    }  
    if (apple == pos) {  
        while (apple == pos) {  
            apple = random_apple(max_x, max_y);  
        }  
        return 1;  
    }  
    return 0;  
}
```


Demo (if time)

- Debug\TwoDime.exe a

(0, 0: Left), -0.559394	(2, 0: Left), 0.0980189	(4, 0: Right), -0.19
(0, 0: Right), 0.0523588	(2, 0: Down), -0.252313	(4, 2: Right), -0.1
(0, 0: Up), 0.0767104	(2, 2: Left), 0.392291	(4, 3: Right), -0.271
(0, 0: Down), -0.174421	(2, 2: Down), 0.0392666	(4, 4: Right), -0.1
(0, 1: Left), -0.480107	(2, 3: Left), 0.00988021	
(0, 1: Right), 0.0223568	(2, 3: Up), 0.198903	
(0, 1: Up), 0.0434685	(2, 4: Up), 0.697903	
(0, 1: Down), 0.0764805	(3, 0: Down), -0.19	
(0, 2: Down), 0.0979209	(3, 1: Left), 0.0999	
(0, 3: Left), -0.1	(3, 1: Down), 0.0099	
(1, 0: Down), -0.569533	(3, 2: Left), 0.0099	
(1, 1: Right), 0.00959766	(3, 2: Right), 0.0989055	
(1, 1: Up), 0.00970387	(3, 2: Up), 0.0186749	
(1, 1: Down), 0.0968491	(3, 4: Up), 0.1	
(1, 4: Right), 0.0296703		
(1, 4: Up), 0.2997		

Massive state space

- 5 by 5
 - plus bag edges
 - possible O or + anywhere
 - will slowly forget last apples place
- In fact, I can't tell where the apples were from the Q-table!
- Does the position matter?
- Character up, down, left, right matters
 - And maybe direction to food
 - Penalise if the agent moves further away

A new lookup

```
struct CharState
{
    char cu;
    char cd;
    char cl;
    char cr;
    int horizontal; // to food -1, 0, +1
    int vertical; // to food -1, 0, +1
    auto operator<=>(const CharState&) const = default;
};
using LookupCharOnly =
    std::map<std::tuple<CharState, Action>, double>;
```

And a snake game (finally!)

```
using Snake = std::deque<Position>;
class Game;
private:
    int max_x{}; int max_y{};
    Snake snake_{ {max_x / 2, max_y / 2} };
    Position apple_{respawn_apple()};
    std::function<Position()> spawn_apple;
    void respawn_apple() {
        do {
            apple = spawn_apple();
        }
        while (std::ranges::find(snake_, apple)
            != snake_.end());
    }
    bool game_over_{ false };

public:
    Game(int max_x, int max_y,
        std::function<Position()> fn);

    [[nodiscard]] bool move(Direction dir);
    bool game_over() const { return game_over_; }
    void reset(Snake snake);
```

Move snake:

```
using Snake = std::deque<Position>;
bool SnakeLib::Game::move(Direction dir) {
    if (!SnakeLib::move(snake_, dir, max_x, max_y)) {
        game_over_ = true;
    }
    else if (eat(snake_.front(), apple)) {
        apple = respawn_apple();
        return true;
    }
    else {
        snake_.pop_back();
    }
    return false;
}
```

And a new environment with a game...

```
class Environment
{
    SnakeLib::Game game;
    SnakeLib::Snake start_snake{game.Snake()};
public:
    Environment(const SnakeLib::Game& game) : game(game) {};
    bool game_over() const { return game.game_over(); }
    void reset() { game.reset(start_snake); };
    FoodDirection food_direction() const;
    int step(Action action);
    CharState state() const;
    SnakeLib::Game& Game() { return game; } // leaky but hey
};
```

... rewards

```
int Environment::step(Action action)
```

- -2 for game over (including out of bag)
- +2 for ate apple
- -1 for further from food
- +1 for nearer food

(Some) Stats

(, , , , [0, -1]: Left), -0.161781

(, , , , [0, -1]: Right), -0.1

(, , , , [0, -1]: Up), 1.27213

(, , , , [0, -1]: Down), 7.40003

(, , , , [0, 1]: Left), 2.56456

(, , , , [0, 1]: Right), 0.466925

(, , , , [0, 1]: Up), 7.22619

(, , , , [1, -1]: Left), 2.1945

(, , , , [1, -1]: Right), 3.04448

(, , , , [1, -1]: Up), 2.35066

(, , , , [1, -1]: Down), 7.68758

(, , , , [1, 0]: Left), -0.1

(, , , , [1, 0]: Right), 7.68254 //...

(, , , O, [1, 0]: Left), 2.18344

(, , , O, [1, 0]: Right), 7.42734

(, , , O, [1, 0]: Up), 1.6151

(, , , O, [1, 0]: Down), 0.0784815

(, , , |, [-1, -1]: Left), 0.665528

(, , , |, [-1, 0]: Up), -0.0901

(, , , |, [-1, 1]: Left), 2.59265

(, , , |, [-1, 1]: Down), -0.1

(, , , |, [0, 1]: Left), 0.37943

(, , , |, [0, 1]: Right), -0.153379

(, , , |, [0, 1]: Up), 6.61384

(, , , |, [0, 1]: Down), 0.081203

(, , +, , [-1, -1]: Left), 0.276949

(, , +, , [-1, -1]: Right), -0.1

(, , +, , [-1, -1]: Up), 0.0830456

(, , +, , [-1, -1]: Down), 7.50895

(, , +, , [-1, 0]: Left), 3.72608

(, , +, , [-1, 0]: Right), 0.052774

(, , +, , [-1, 0]: Down), -0.126848

(, , +, , [-1, 1]: Left), -0.0901

(, , +, , [-1, 1]: Right), 0.27479

(, , +, , [-1, 1]: Up), 8.14921

(, , +, , [-1, 1]: Down), -0.228743

Demo

- Maybe time for a demo? Use Apples.exe q_table_1000.txt

```
|  +*+ |  
| O  ++ |  
|    ++ |  
|    ++ |  
|    ++ |
```

Won 0, lost 1


Scores

20

SNAKE!!!!!!!!!!

- Configure grid size
 - Started 5 by 5,
 - Then try 50 by 45
- Save the q_table
 - So we can read in a trained version
 - And use this to play Snake

Score 30



Demo


- Debug\DemoTrainedRL.exe ..\q_table_large_10000.txt
- Also a short screen capture here:
https://www.linkedin.com/posts/francesbuontempo_i-have-attempted-to-teach-my-machine-how-activity-7303789398104395776-CX8n/

Summary


- Reinforcement learning is one type of “AI”
- It doesn’t need labelled data
 - Simple versions like Q-learning store data in tables
 - A kind of Markov decision process, mapping states to actions
 - <https://accuconference.org/2025/session/a-very-small-language-model> (Jez’s talk)
- I used a temporal difference Q-learning approach
 - There are other approaches
- It can discover how to play games
 - Can be used in the “fine tuning” of LLMs
 - And do more useful stuff
 - Compiler optimization, logistics, power grid management, video compression

Other things to go play with


- <https://play.battlesnake.com/>
- https://compiler gym.com/getting_started.html
 - a toolkit for applying reinforcement learning to compiler optimization tasks.
- OpenQI's Gym: <https://gymnasium.farama.org/index.html>
- Arcade Learning Environment (ALE): <https://ale.farama.org/environments/>
 - Asteroids, PacMan...
 - Env can be RGB image
- Deep RL
 - Started the Atari games learning
 - Uses a convolutional neural network instead of a Q-table.
 - Playing Atari with Deep Reinforcement Learning, 2013, <https://arxiv.org/abs/1312.5602>
- Deep Reinforcement Learning in Pac-man
 - <https://www.youtube.com/watch?v=QilHGSYbjDQ>



Introducing C++



Learn C++ by Example



Genetic Algorithms and
Machine Learning for
Programmers

<https://www.oreilly.com/library/view/introducing-c/9781098178130/>

<https://mng.bz/1aVV>

<https://pragprog.com/titles/fbmach/genetic-algorithms-and-machine-learning-for-programmers/>