

# Building the Backbone for Market Data in Modern C++

2025

**Jody Hagins,  
Robert Leahy  
& Gianluca Delfino**

# Building the Backbone for Market Data in Modern C++

Jody Hagins – Director, Software Engineering  
[jody.hagins@lsegu.com](mailto:jody.hagins@lsegu.com)  
[coachhagins@gmail.com](mailto:coachhagins@gmail.com)

Gianluca Delfino – Director, Software Engineering  
[gianluca.delfino@lsegu.com](mailto:gianluca.delfino@lsegu.com)  
[gianluca.delfino@gmail.com](mailto:gianluca.delfino@gmail.com)

Robert Leahy – Lead Software Engineer  
[robert.leahy@lsegu.com](mailto:robert.leahy@lsegu.com)  
[rleahy@rleahy.ca](mailto:rleahy@rleahy.ca)



**LSEG**

# What we do in the “Real Time **Ultra** Direct” Team

...and how we do it!



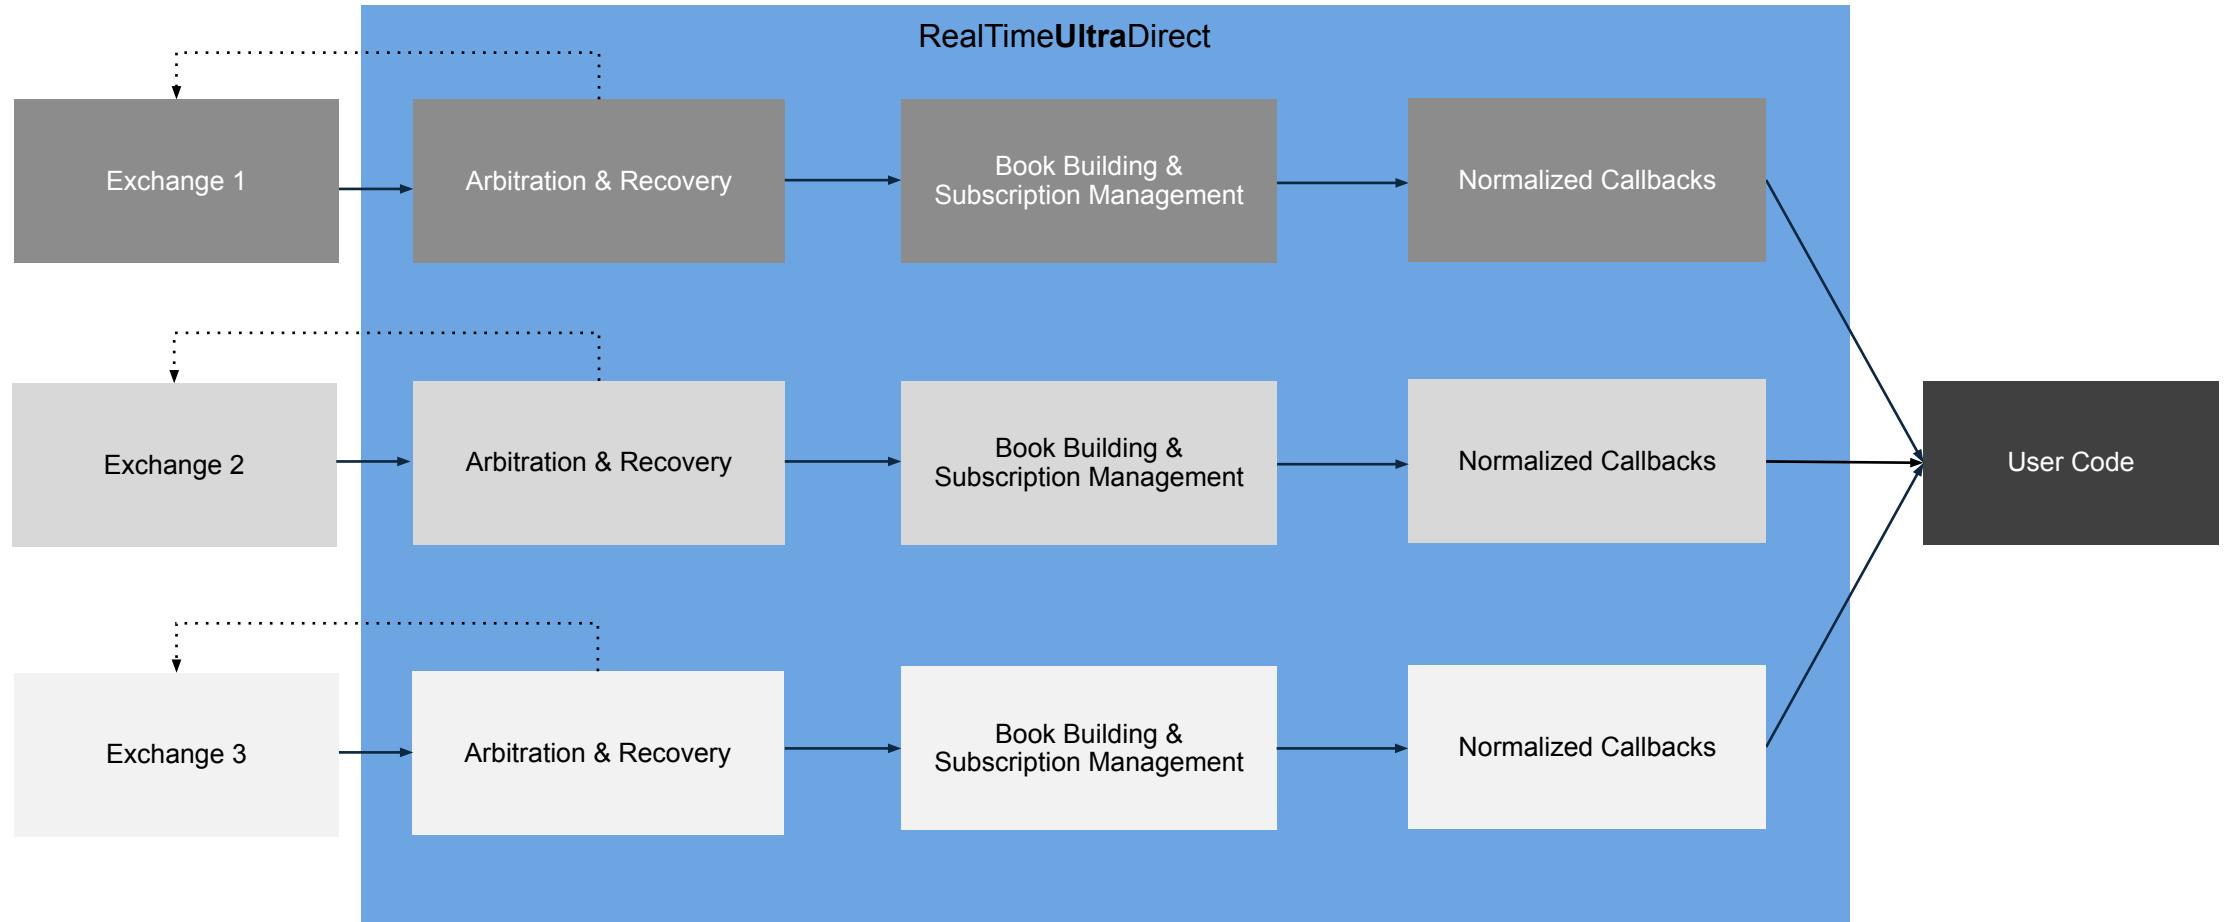
**LSEG**

## What is RealTime**Ultra**Direct

---

- C++ library that can handle millions of messages per second
- Mission: democratize HFT
- We provide *Normalization* for hundreds of feeds over many exchanges
- Low barrier of entry, easy to set up and start using.
- Comes with Batteries included:
  - built for every Linux platform/Compiler combination
  - C++/Java bindings
  - Performance analysis tools
- Extensible:
  - Buy single feed-handlers / pay for what you use
- “Extra mile” Support:
  - Direct access to developers for technical issues

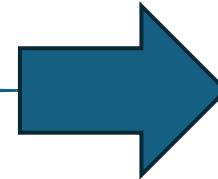
# One “Normalized” Interface to rule them ALL



# One “Normalized” Interface to rule them ALL

## example “RAW”

- ComplexInstrumentDefinition
- SingleSideUpdate
- TwoSideUpdate
- TopTrade
- AddOrderLong
- AddOrderShort
- AddOrderExpanded
- OrderExecuted
- ModifyOrderLong
- ModifyOrderShort
- DeleteOrder
- TradeLong
- TradingStatus
- AuctionUpdate
- AuctionCancel
- AuctionTrade
- ...



## Normalize d

- ProductAnnouncement
- ProductStatistics
- AddOrder
- ModifyOrder
- CancelOrder
- PriceLevelUpdate
- OrderImbalance
- AuctionSummary
- ProductStatus
- FeedStatus
- MarketStatus
- BBOQuote
- NBBOQuote
- Trade
- TradeCorrection
- IndexUpdate
- ...



# How do we do it? (API)



## A good API is shaped like an Onion

- The external layer is *Normalized Messages* Callbacks

### Normalization Api:

```
void OnProductAnnouncement(  
    const feed::ProductAnnouncement& announcement)
```



# How do we do it? (API)



## A good API is shaped like an Onion

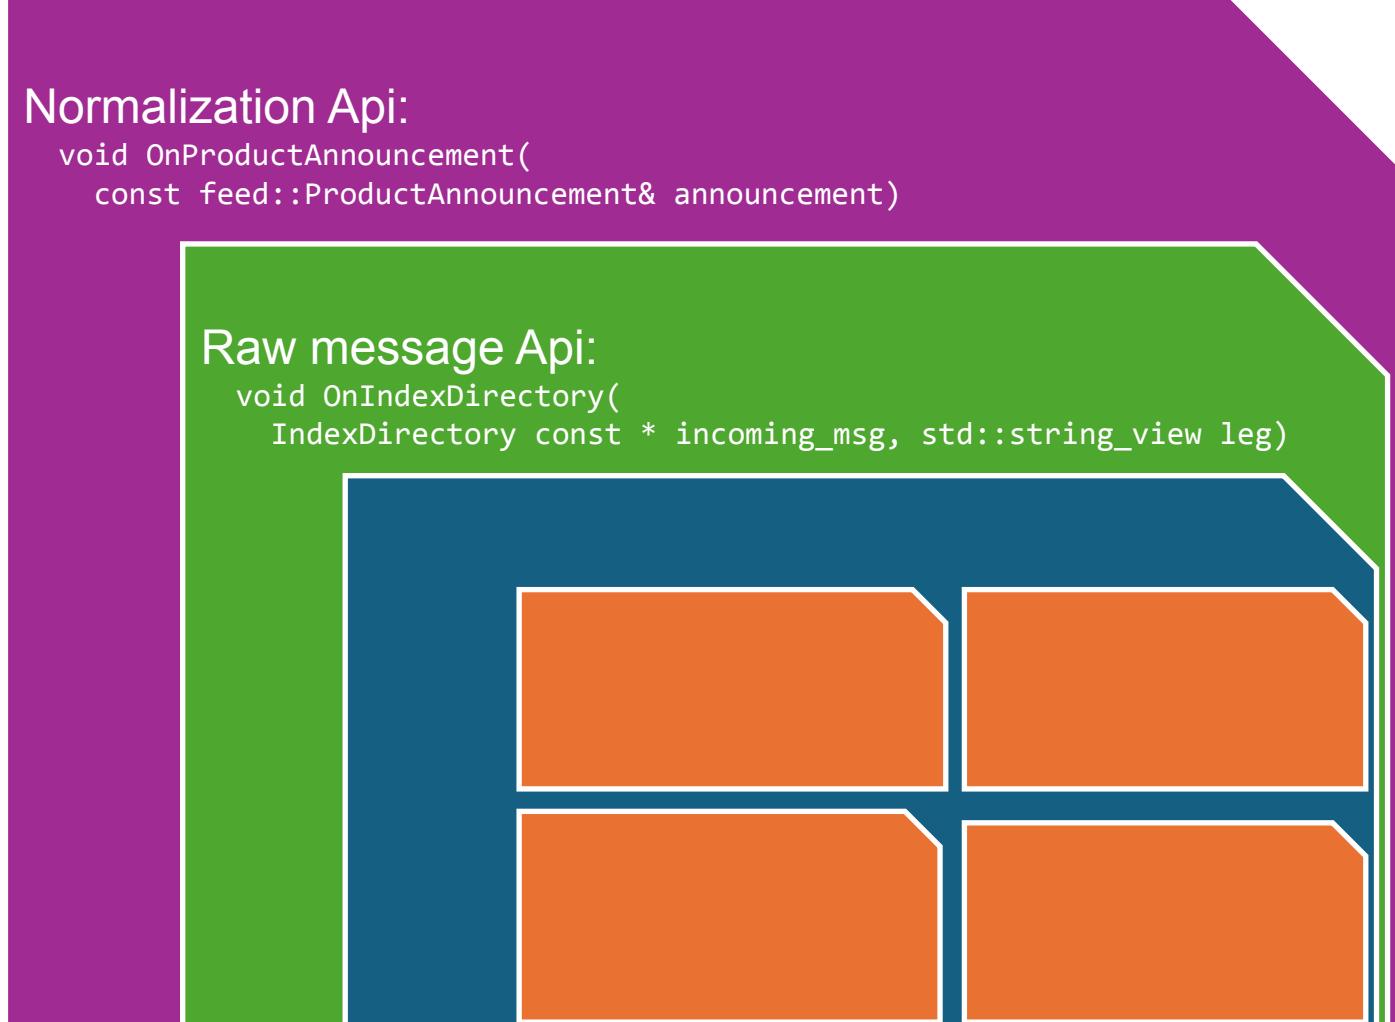
- The external layer is *Normalized Messages* Callbacks
- Optionally users can access “Raw” Messages

### Normalization Api:

```
void OnProductAnnouncement(  
    const feed::ProductAnnouncement& announcement)
```

### Raw message Api:

```
void OnIndexDirectory(  
    IndexDirectory const * incoming_msg, std::string_view leg)
```



# How do we do it? (API)



## A good API is shaped like an Onion

- The external layer is *Normalized Messages* Callbacks
- Optionally users can access “Raw” Messages

### Normalization Api:

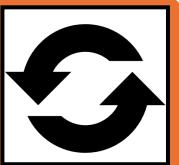
```
void OnProductAnnouncement(  
    const feed::ProductAnnouncement& announcement)
```

### Raw message Api:

```
void OnIndexDirectory(  
    IndexDirectory const * incoming_msg, std::string_view leg)
```

### Internal APIs:

Arbitration & Recovery



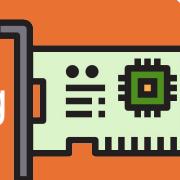
Book Building



Pcap Handling



Device Handling

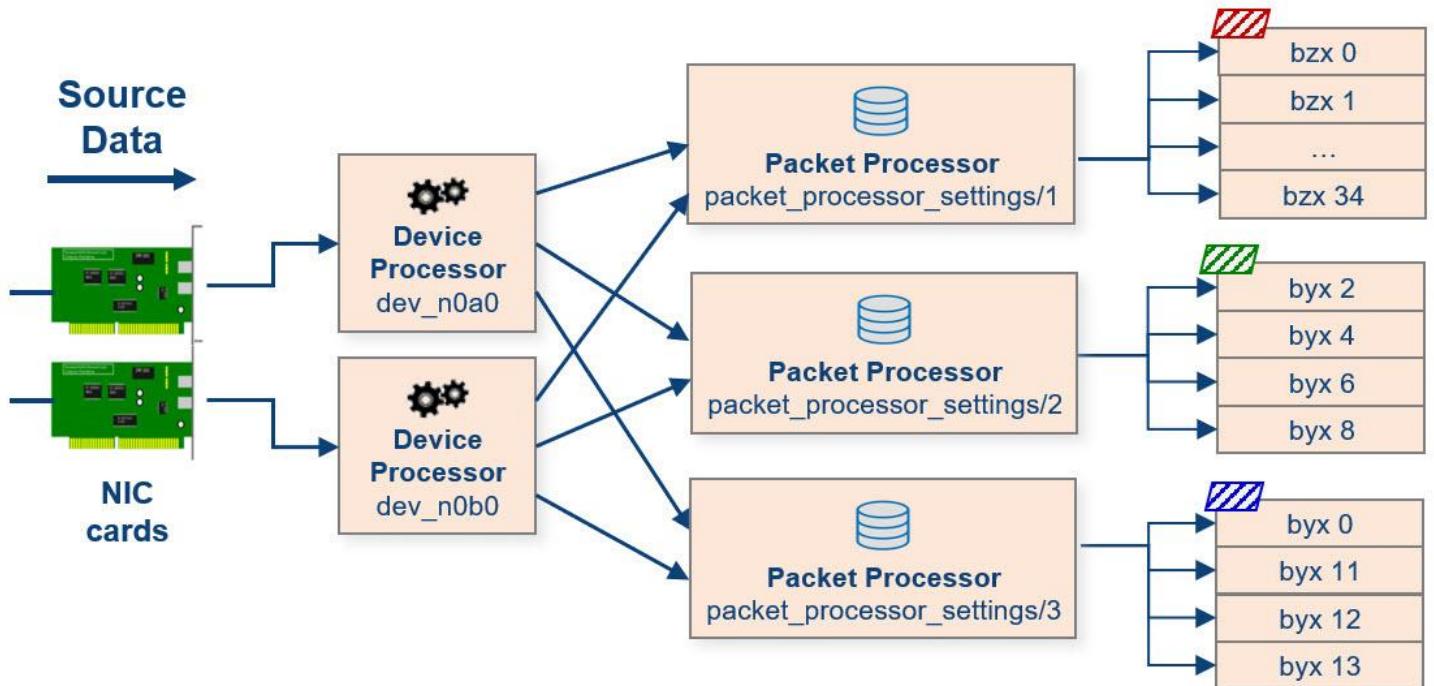


# How do we do it? (API)

## Fully Configurable:

- Devices support `ef_vi` kernel bypass
- Packet processed on specified *threads*

## Example: Feed Breakout



# How do we do it? (API)

---

## Our Tech Stack

- More than 2.5 Million lines of C++
- Java bindings are not included in this breakdown
- C++17

```
$ cloc .
 15270 text files.
 14398 unique files.
 928 files ignored.
```

[github.com/AlDanial/cloc v 2.04 T=51.72 s \(278.4 files/s, 85414.0 lines/s\)](https://github.com/AlDanial/cloc)

Language	files	blank	comment	code
C++	6313	192258	113575	1654211
C/C++ Header	6296	158230	320269	962718
XML	379	9548	7466	453605
JSON	80	0	0	303255
Text	63	13468	0	103375
CSV	15	1021	0	44702
CMake	937	3946	6215	22415
Python	81	2315	610	10601
XSD	17	151	55	6213
JavaScript	18	318	768	6122
TypeScript	40	768	1329	5490
Markdown	48	852	1	2663
Bourne Shell	28	281	113	1384
PHP	30	193	398	1005
SCSS	7	49	39	996
Pascal	14	191	999	840
CSS	6	78	33	717
SVG	1	0	0	685
YAML	3	7	14	634
HTML	11	16	29	310
Dockerfile	5	53	30	180
vim script	1	14	8	45
diff	2	8	25	37
Visual Studio Solution	1	1	1	18
C	1	1	0	3
INI	1	0	0	2
SUM:	14398	383767	451977	3582226

# How do we do it? (CI)

---

## Our CI for every commit:

- clang18-ccache
- gcc-14-ccache
- devtoolset-7-ccache

## Overnight:

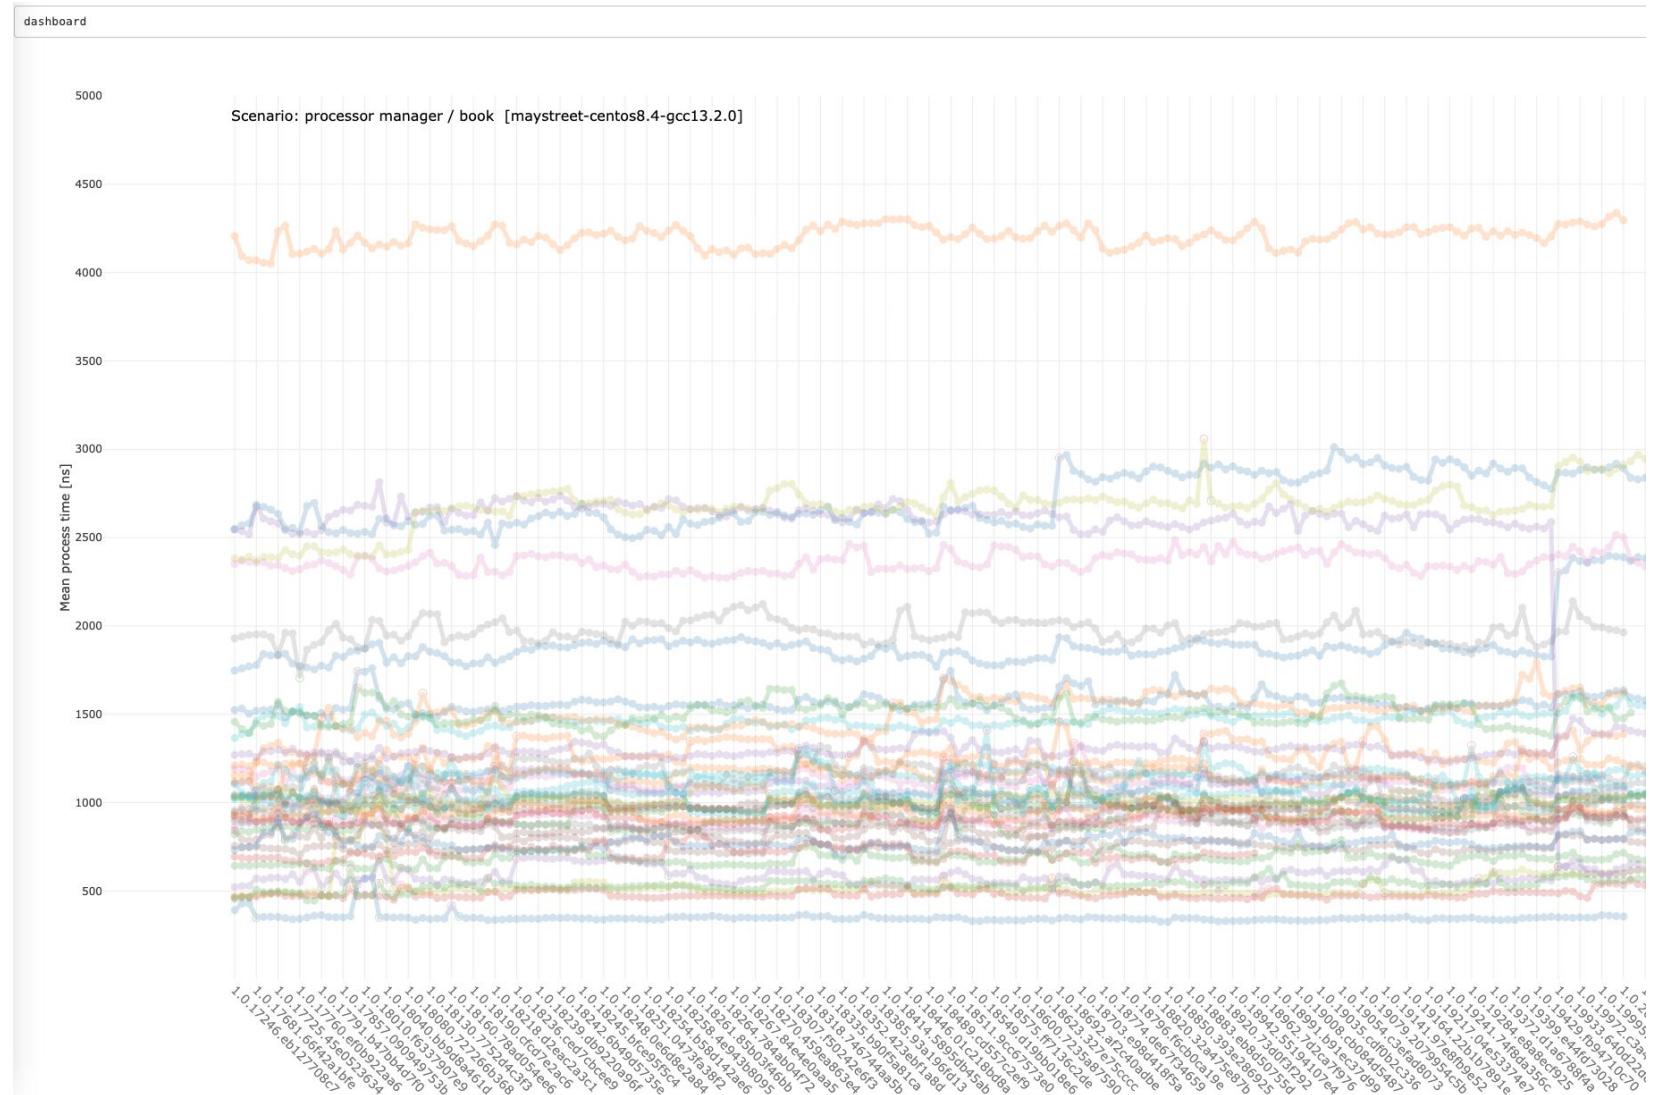
- msan
- tsan
- ubsan
- asan-lsan
- clang-tidy
- memcheck



# Performance Monitoring

## Continuous monitoring

- Every build is tested with a set of thorough integration tests on production hardware.

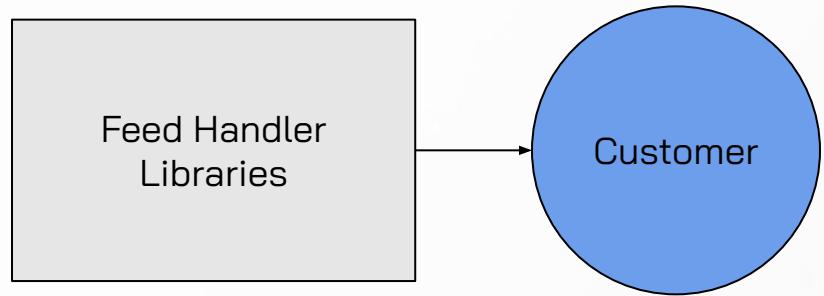


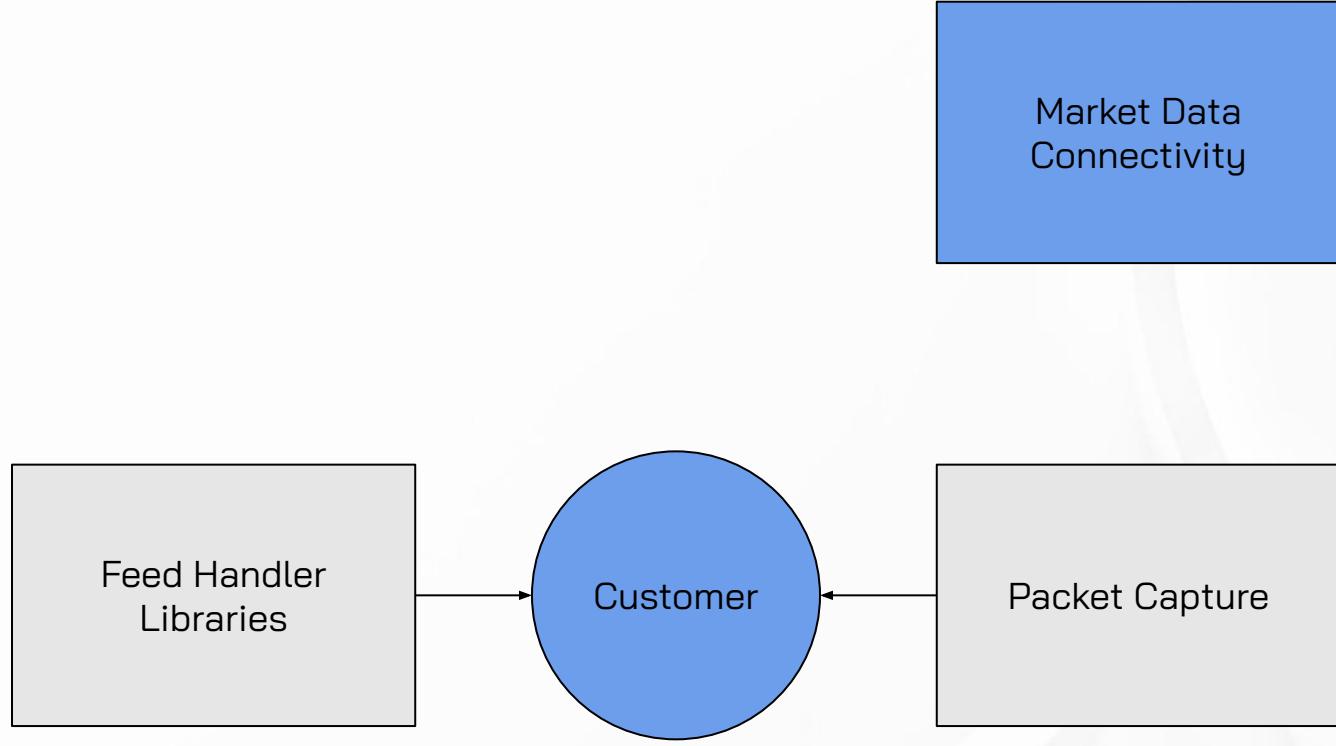
## What is next

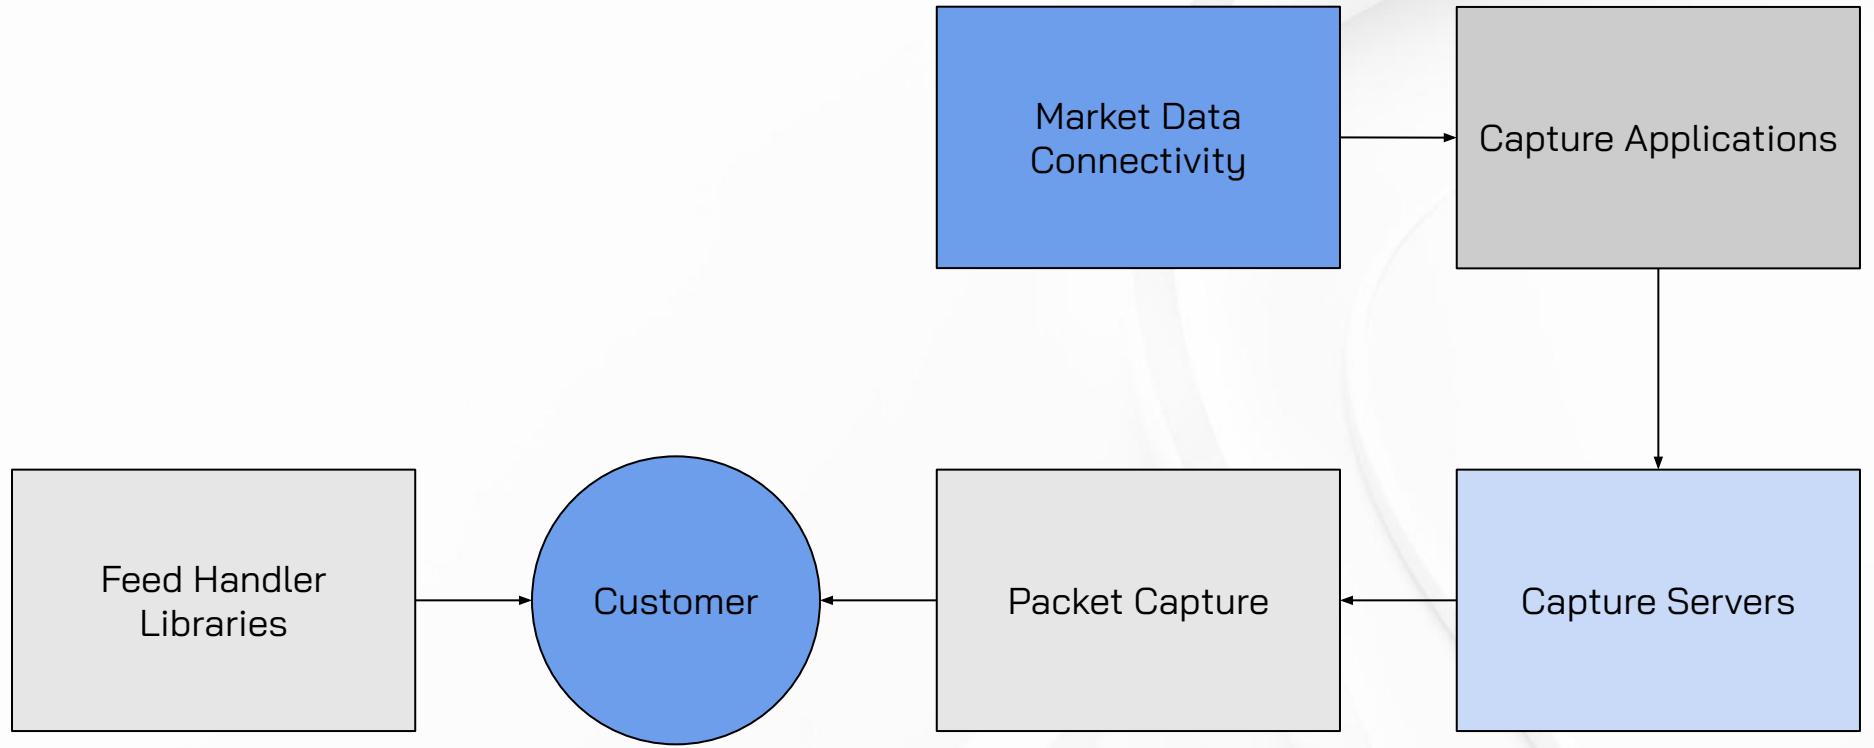
---

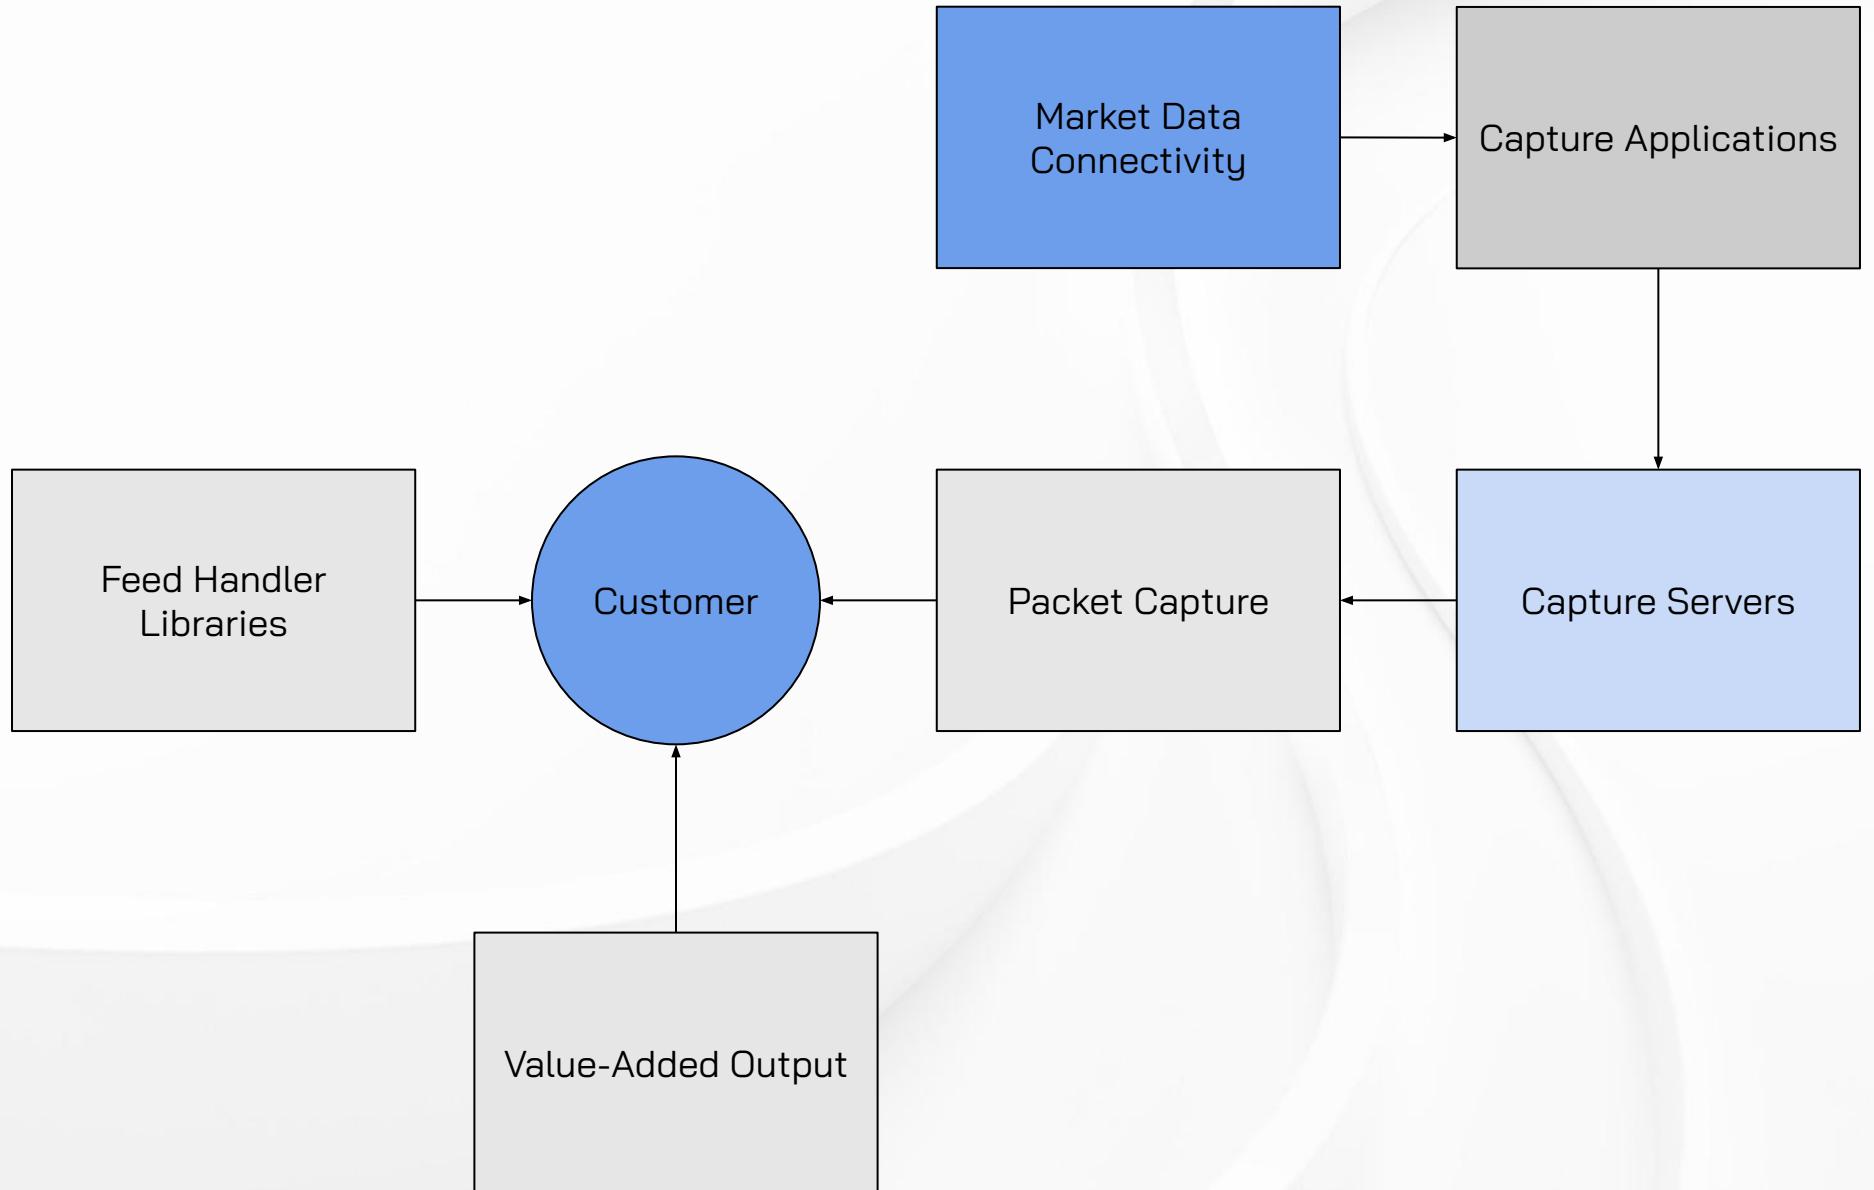
- Better automatic Normalization Code generation
- Customization Options for Normalization (different “flavours”)
- C++20/23 (**Concepts!**)
- Moar faster

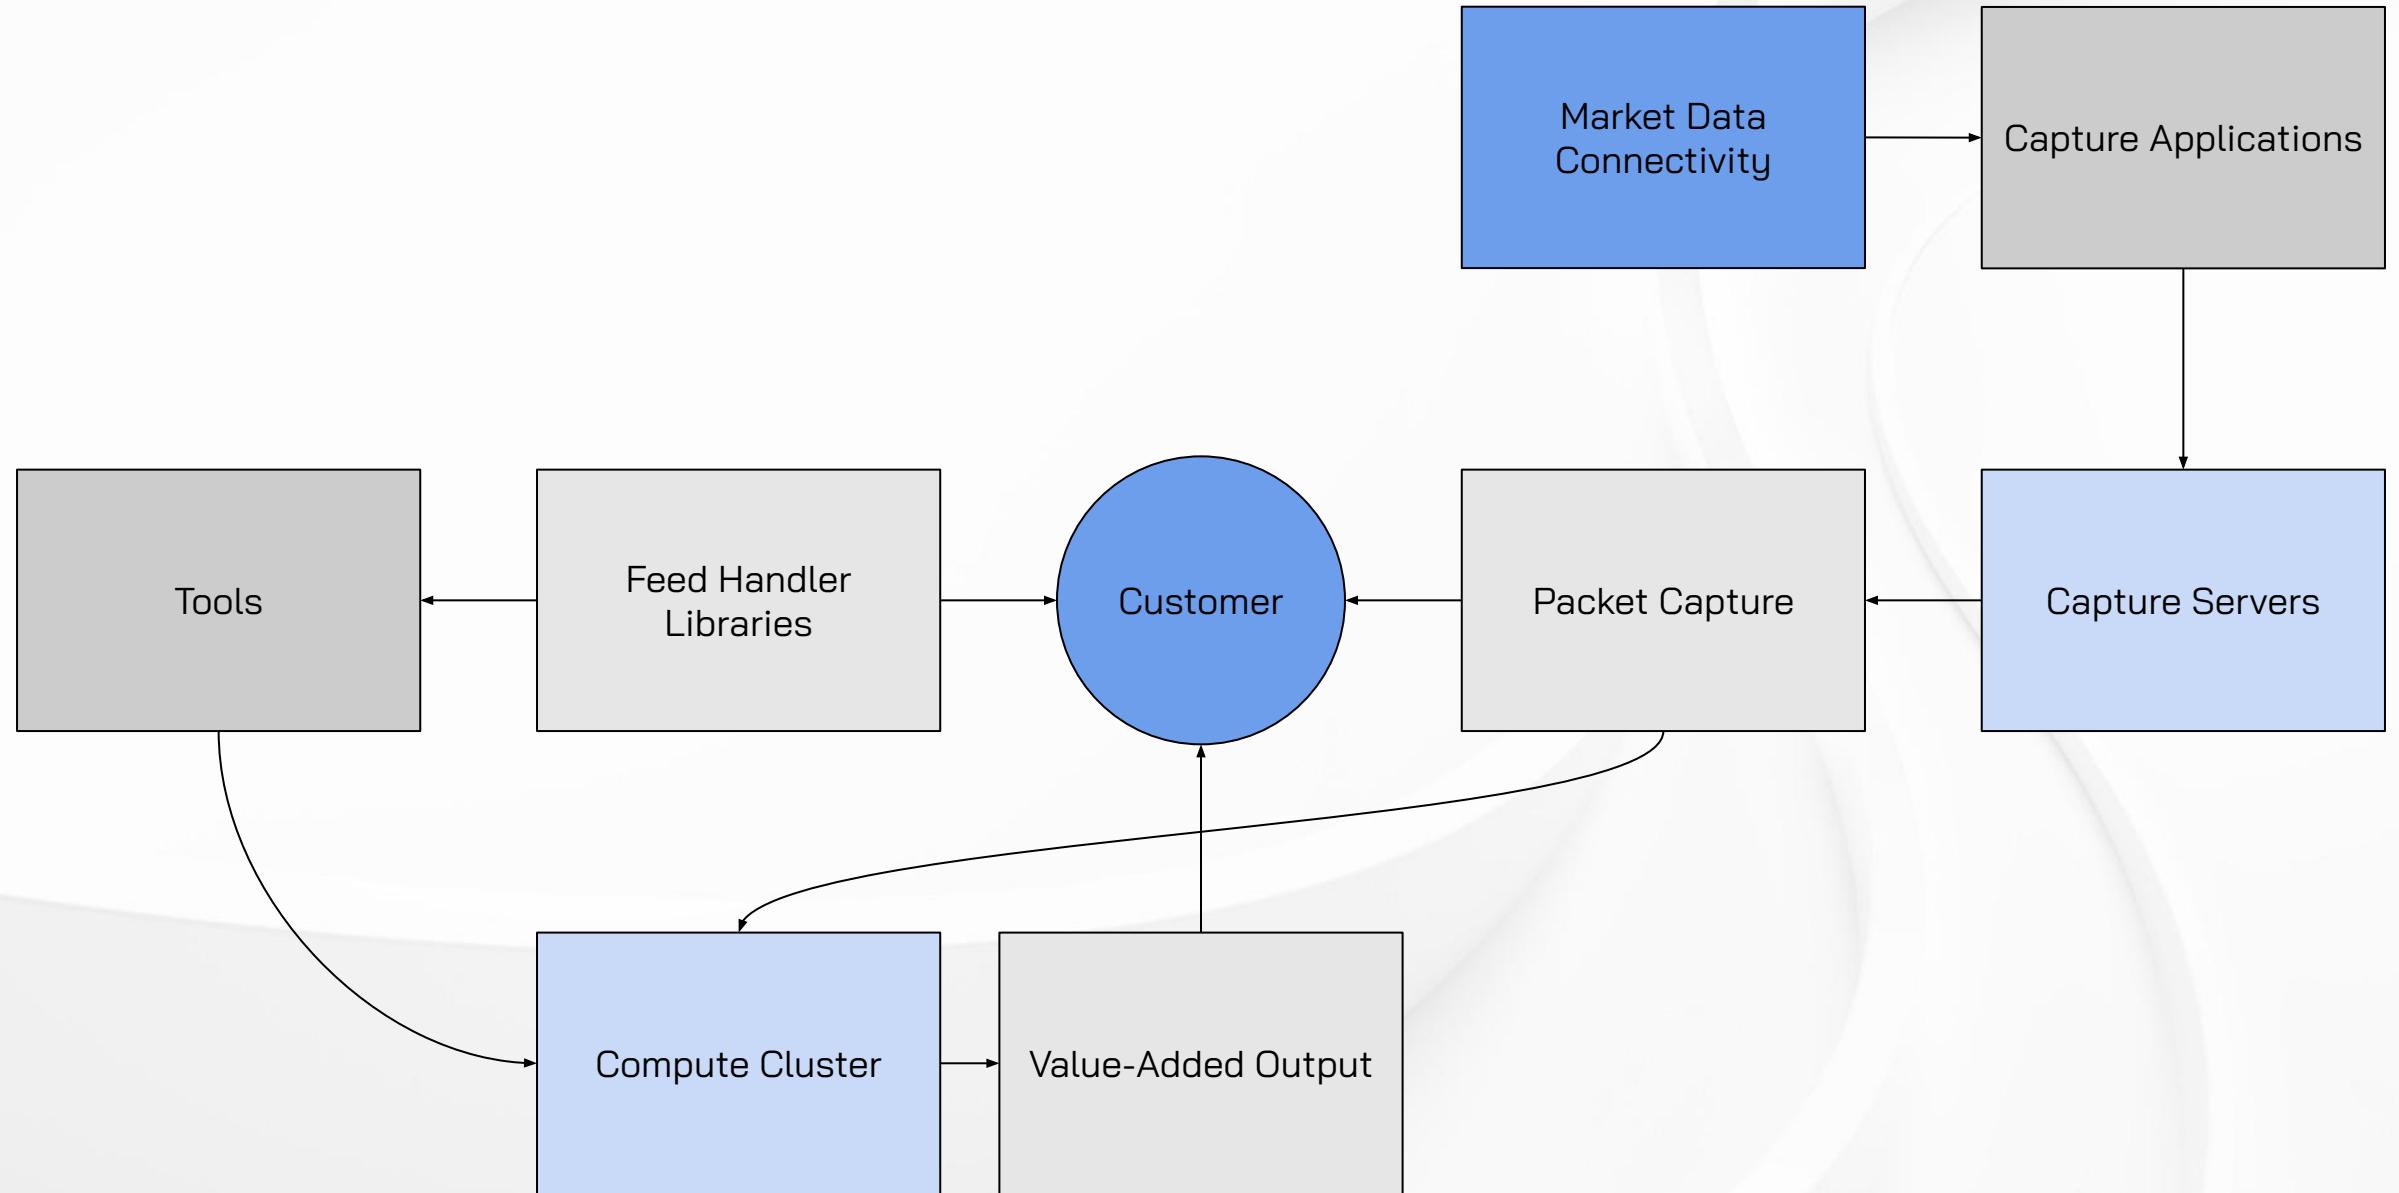
# Tools Team

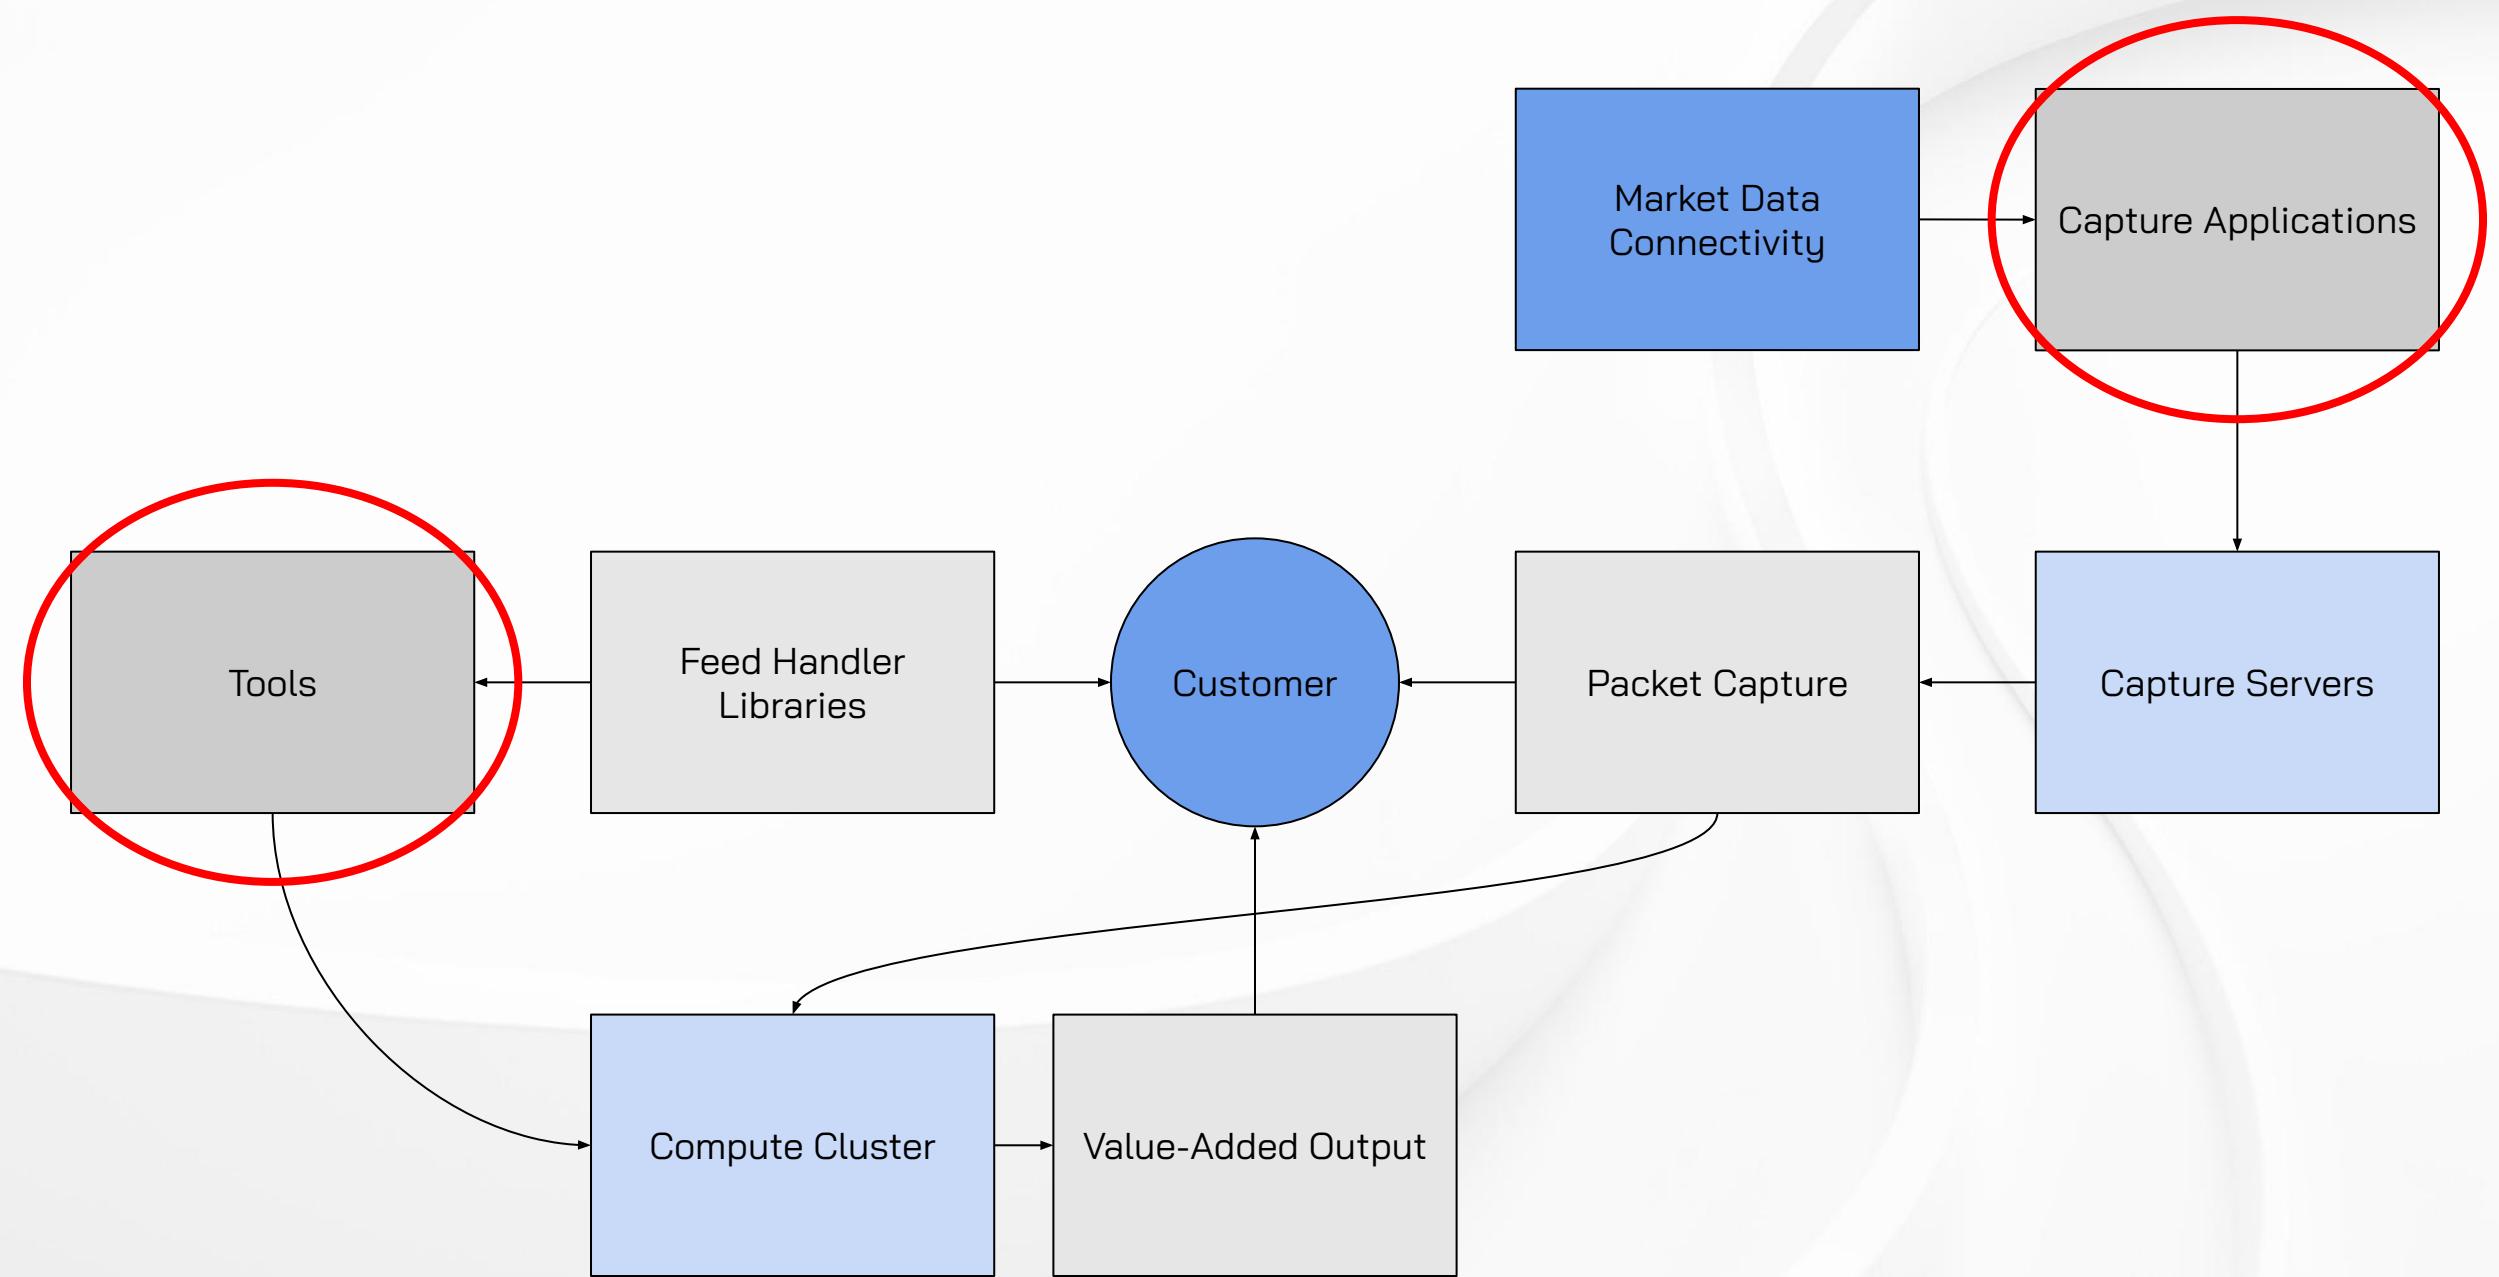












# Breaking it Down

## What're the projects behind the offerings?

### Capture

`mstcaprecv`

Captures raw data from network using kernel bypass.

`mstcapwrite`

Writes raw data to raw capture files.

`mstcapmon`

Generates quality metrics about captured data.

`mstsplit`

Post-processes raw capture into client-facing capture offerings.

`replay`

Replays pcaps onto the network for functional testing and simulation.

# Breaking it Down

## What're the projects behind the offerings?

### Capture

`mstcaprecv`

Captures raw data from network using kernel bypass.

`mstcapwrite`

Writes raw data to raw capture files.

`mstcapmon`

Generates quality metrics about captured data.

`mstsplit`

Post-processes raw capture into client-facing capture offerings.

`replay`

Replays pcaps onto the network for functional testing and simulation.

### Normalization

`mstnorm`

Transforms market data to various formats (CSV, JSON, Parquet, et cetera) with or without additional aggregate measures.

`mstanalyze`

Pluggable, general purpose processing pipeline.

# Breaking it Down

## What're the projects behind the offerings?

### Capture

`mstcaprecv`

Captures raw data from network using kernel bypass.

`mstcapwrite`

Writes raw data to raw capture files.

`mstcapmon`

Generates quality metrics about captured data.

`mstsplit`

Post-processes raw capture into client-facing capture offerings.

`replay`

Replays pcaps onto the network for functional testing and simulation.

### Normalization

`mstnorm`

Transforms market data to various formats (CSV, JSON, Parquet, et cetera) with or without additional aggregate measures.

`mstanalyze`

Pluggable, general purpose processing pipeline.

### MDX

Proprietary, distributed, time series database engine.

### High Performance Query

REST- and WebSocket-accessible interface for querying market data built on top of MDX.



US 20220044259A1

(19) **United States**

(12) **Patent Application Publication**

**Douglas et al.**

(10) **Pub. No.: US 2022/0044259 A1**

(43) **Pub. Date:** **Feb. 10, 2022**

---

(54) **DATABASE MANAGEMENT SYSTEMS AND METHODS USING DATA NORMALIZATION AND DEFRAGMENTATION TECHNIQUES**

(71) Applicant: **Maystreet, Inc.**, New York, NY (US)

(72) Inventors: **Niall Douglas**, Kerry Pike (IE); **Robert Leahy**, New York, NY (US); **Michael Lehr**, New York, NY (US)

(73) Assignee: **Maystreet, Inc.**, New York, NY (US)

(21) Appl. No.: **16/986,646**

(22) Filed: **Aug. 6, 2020**

(52) **U.S. Cl.**

CPC ..... **G06Q 30/0201** (2013.01); **G06F 16/1724** (2019.01); **G06F 16/128** (2019.01); **G06F 16/178** (2019.01); **G06F 16/113** (2019.01)

(57)

### **ABSTRACT**

Improved systems and methods for database management using data normalization and defragmentation techniques are provided. At least one exchange processor in communication with an exchange computer system receives market data from the exchange computer system, processes the market information, and transmits the market data to a master processor. The master processor receives the market data, processes the data using at least one normalization process to generate normalized data including an intra-day file and an archival file, and stores the intra-day file and the

# Breaking it Down

## What're the projects behind the offerings?

### Capture

`mstcaprecv`

Captures raw data from network using kernel bypass.

`mstcapwrite`

Writes raw data to raw capture files.

`mstcapmon`

Generates quality metrics about captured data.

`mstsplit`

Post-processes raw capture into client-facing capture offerings.

`replay`

Replays pcaps onto the network for functional testing and simulation.

### Normalization

`mstnorm`

Transforms market data to various formats (CSV, JSON, Parquet, et cetera) with or without additional aggregate measures.

`mstanalyze`

Pluggable, general purpose processing pipeline.

### MDX

Proprietary, distributed, time series database engine.

### High Performance Query

REST- and WebSocket-accessible interface for querying market data built on top of MDX.

# Breaking it Down

## What're the projects behind the offerings?

### Capture

`mstcaprecv`

Captures raw data from network using kernel bypass.

`mstcapwrite`

Writes raw data to raw capture files.

`mstcapmon`

Generates quality metrics about captured data.

`mstsplit`

Post-processes raw capture into client-facing capture offerings.

`replay`

Replays pcaps onto the network for functional testing and simulation.

C++11

### Normalization

`mstnorm`

Transforms market data to various formats (CSV, JSON, Parquet, et cetera) with or without additional aggregate measures.

`mstanalyze`

Pluggable, general purpose processing pipeline.

### MDX

Proprietary, distributed, time series database engine.

### High Performance Query

REST- and WebSocket-accessible interface for querying market data built on top of MDX.

# Breaking it Down

What're the projects behind the offerings?

## Capture

`mstcaprecv`

Captures raw data from network using kernel bypass.

`mstcapwrite`

Writes raw data to raw capture files.

`mstcapmon`

Generates quality metrics about captured data.

`mstsplit`

Post-processes raw capture into client-facing capture offerings.

`replay`

Replays pcaps onto the network for functional testing and simulation.



## Normalization

`mstnorm`

Transforms market data to various formats (CSV, JSON, Parquet, et cetera) with or without additional aggregate measures.

`mstanalyze`

Pluggable, general purpose processing pipeline

**C++20**

## MDX

Proprietary distributed, time series database engine.

## High Performance Query

REST- and WebSocket-accessible interface for querying market data built on top of MDX.

**C++20**

# Breaking it Down

What're the projects behind the offerings?

## Capture

`mstcaprecv`

Captures raw data from network using kernel bypass.

`mstcapwrite`

Writes raw data to raw capture files.

`mstcapmon`

Generates quality metrics about captured data.

`mstsplit`

Post-processes raw capture into client-facing capture offerings.

`replay`

Replays pcaps onto the network for functional testing and simulation.



## Normalization

`mstnorm`

Transforms market data to various formats (CSV, ISO1, Parquet, et cetera) with or without additional aggregate measures.

`mstanalyze`

Pluggable, general purpose processing pipeline.

**C++23**

**C++20**

## MDX

Proprietary distributed, time series database engine.

## High Performance Query

REST- and WebSocket-accessible interface for querying market data built on top of MDX.

**C++20**

## Deploying the Networking TS



Failing Successfully:  
Reporting and Handling  
Errors

ROBERT LEAHY

20  
21 | October 24-29

+ 21



Taking a Byte Out of C++  
Avoiding Punning by Starting Lifetimes

ROBERT LEAHY

20  
21 | October 24-29

+ 22



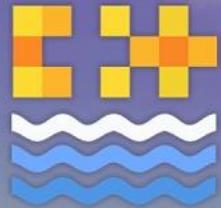
ROBERT LEAHY

20  
22 | September 12th-16th

+

21

+



C++ now 2024  
Aspen, Colorado  
CppNow.org

Video Sponsorship Provided By  
millennium  
think-cell



Let  $e$  be the expression  $\text{do-run}(s)$ , and the type  $E$  be  $\text{decltype}((e))$ . Then  $\text{run}(s)$  is expression-equivalent to:

- $e$ , if that expression is well-formed and  $E$  is of the form  $\text{outcome}::\text{result}(\text{std}::\text{variant}(<\!\!>))$ , otherwise
- $\text{outcome}::\text{result}(\text{std}::\text{variant}(<\!\!>))(e)$  if that expression is well-formed and  $E$  is of the form  $\text{std}::\text{variant}(<\!\!>)$ , otherwise
- $\text{outcome}::\text{result}(\text{std}::\text{variant}(<\!\!>))(e.\text{assume\_value}())$  if that expression is well-formed,  $E$  is of the form  $\text{outcome}::\text{result}(<\!\!>)$ , and  $\text{bool}(e)$  is true, otherwise
- $\text{outcome}::\text{result}(\text{std}::\text{variant}(<\!\!>))(e.\text{as\_failure}())$  if that expression is well-formed, and  $E$  is of the form  $\text{outcome}::\text{result}(<\!\!>)$ , otherwise
- $\text{outcome}::\text{result}(\text{std}::\text{variant}(<\!\!>))(e)$  if that expression is well-formed.

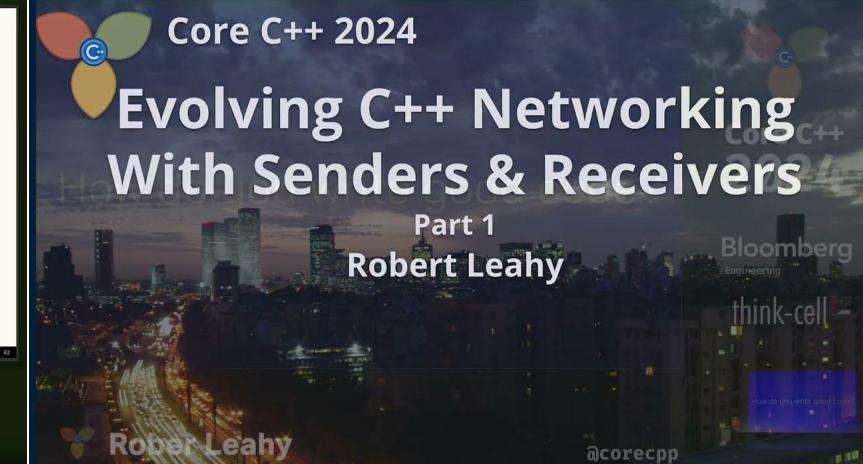
Otherwise,  $\text{run}(s)$  is ill-formed.

An Adventure in Modern Library Design

Robert Leahy

## Fantastic Bugs and How to Test Them

Robert Leahy





# Come to my lightning talk!



# When Do You Know connect Doesn't Throw?

Document Number: P3388R2

Date: 2025-04-01

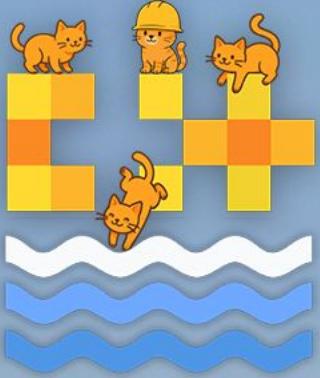
Reply-to: Robert Leahy <rleahy@rleahy.ca>

Audience: LWG

## Abstract

This paper proposes a change which will enable earlier determination that connecting a certain sender with a certain receiver will never throw an exception.

## Background



# Extending std::execution Further

Higher-Order Senders and the Shape of Asynchronous Programs

**Robert Leahy**

**2025**

# A Sentinel for Null-Terminated Strings

Document #: P3705R0  
Date: 2025-05-17  
Project: Programming Language C++  
Audience: SG-9 Ranges  
LEWG  
Reply-to: Eddie Nolan  
[<eddiejnolan@gmail.com>](mailto:eddiejnolan@gmail.com)

## Contents

- 1 Motivation
- 2 Design
  - 2.1 exposition-only helper concept *default-initializable-and-equality-comparable-iter-value*
  - 2.2 `null_sentinel` pass-by-value design
    - 2.2.1 Null sentinel
    - 2.2.2 Operations
  - 2.3 `null_sentinel` pass-by-reference design
    - 2.3.1 Null sentinel
    - 2.3.2 Operations

# **std::constant\_wrapper**

Document #: P2781R8  
Date: 2025-03-15  
Project: Programming Language C++  
Audience: LEWG  
LWG  
Reply-to: Hana Dusíková  
[<hanicka@hanicka.net>](mailto:<hanicka@hanicka.net>)  
Matthias Kretz  
[<m.kretz@gsi.de>](mailto:<m.kretz@gsi.de>)  
Zach Laine  
[<whatwasthataddress@gmail.com>](mailto:<whatwasthataddress@gmail.com>)

## **Contents**

- [1 Changelog](#)
  - [1.1 Changes since R0](#)
  - [1.2 Changes since R1](#)
  - [1.3 Changes since R2](#)
  - [1.4 Changes since R3](#)
  - [1.5 Changes since R4](#)



# Summary

## What do we do?

LSEG Data & Analytics...

- ...writes modern C++

- ...improves C++ itself

- ...creates software for modern financial markets

- ...creates and supports services for modern financial market participants

**LSEG DATA &  
ANALYTICS**

# We're Hiring!

Barry Rindenow – Talent Acquisition Partner  
barry.rindenow@lsegu.com

<https://www.lseg.com/en/careers>



**LSEG**