

# Mind the Gap

(Between Your Code and Your Toolchain)

**Yannic Staudt**

**2025**



# What's in here?



# THE DEBUG SYMBOL'S GUIDE TO THE GALAXY

Mostly harmless, occasionally symbol-less.



# DWARF, the Linker, and Everything

- You start to have problems with linking typically when your binary reaches ~2GB
- By default, **gcc** and **clang** assume `-mmodel=small`

*“the program and its symbols must be linked in the lower 2 GB of the address space”*

Exceed that and you get all sorts of *funny* errors

Exacerbated by debug symbols, lesser levels build optimizations, etc....



# Don't Panic – Enter Debug Fission

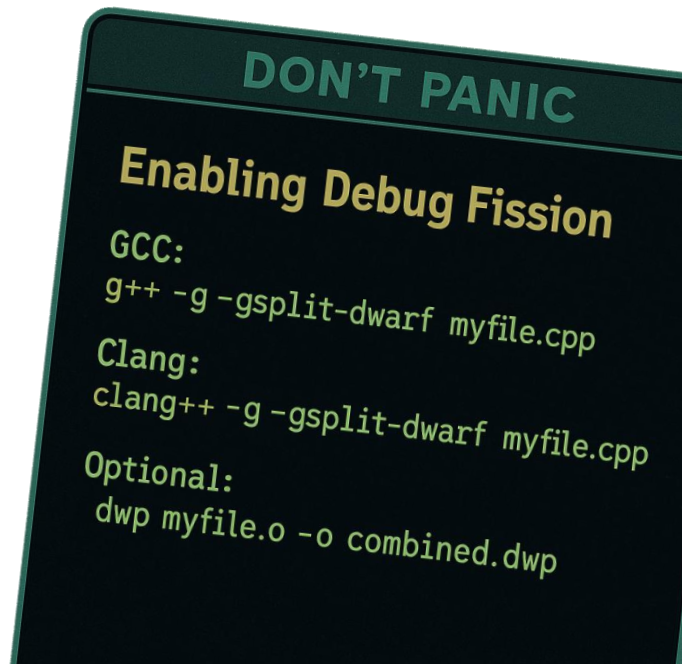
Can be done at compile time – just add `-gsplit-dwarf` to your compiler invocation

DONE!

Even easier handling when you build yourself a debug symbol package (“DWARF package or .dwp”) that contains all the symbols from your programs as you ship it.

If you want a great read-up on that topic:

<https://maskray.me/blog/2022-10-30-distribution-of-debug-information>



# Remember

The answer to “*Why is my linker screaming?*” is rarely 42.

It’s probably just too many symbols.

Use debug fission. **Don’t Panic.**

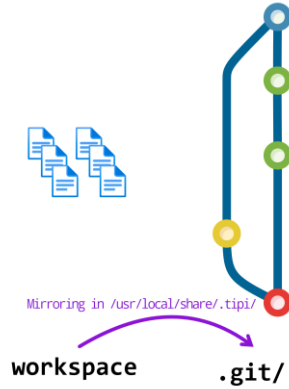
# THE GREAT BRITISH BUILD-OFF

The judges are looking for a light, flaky  
cache with minimal linking.





Caching is  
**tricky**



# Make it invariant

Less variance == higher hit rate & less work



# Controlled environments

Containers are your cache's friend



# Two levels of caching for maximum effectiveness

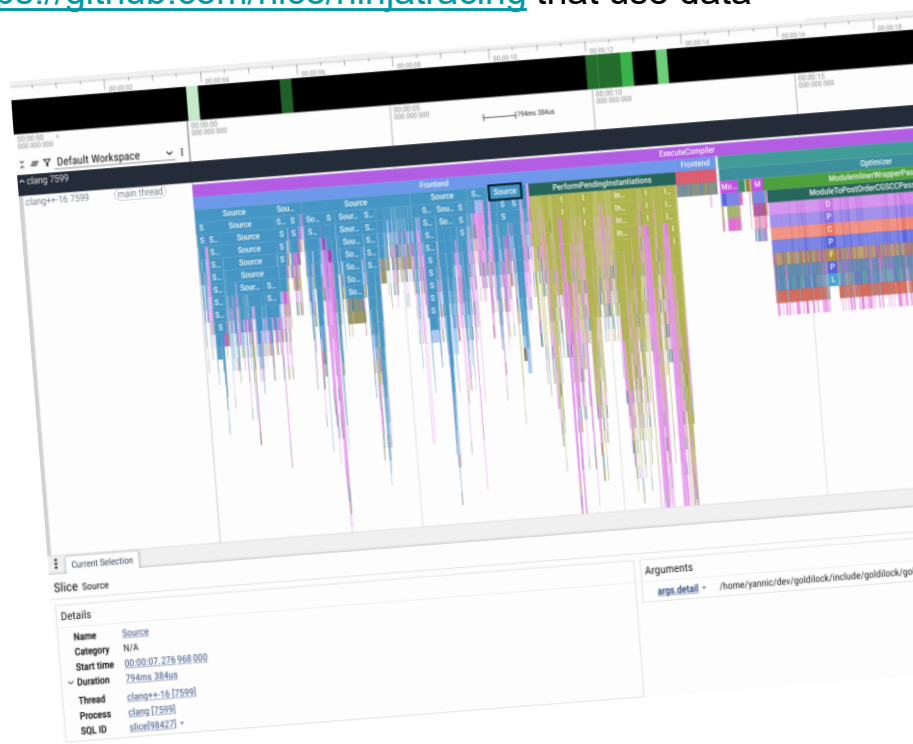
# ONLY TOOLS AND HORSES





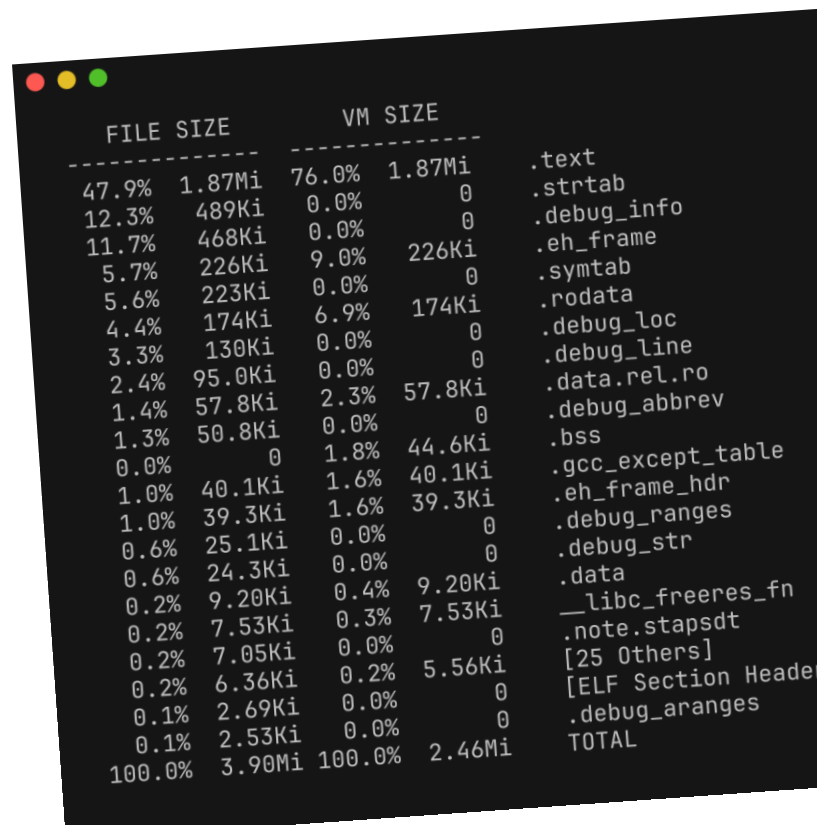
# When did you last profile your build?

- Clang users have it easy: `-ftime-trace=trace.json -ftime-trace-granularity=1`
- If you're using ninja, there's some great tools like <https://github.com/nico/ninjabracing> that use data from ninja's build log
- For MSVC users, look into `vcperf`
- **Soon on our EngFlow platform dashboard!**



# My binary is overweight. Who's been feeding it biscuits??

- Give google's bloaty a try it's **awesome(!)**  
<https://github.com/google/bloaty>
- A size profiler for binaries
- See what contributes most to binary size at the symbol, file, or section level.
- Shows where debug info bloats your binary, or what code/data is causing file size creep.
- Supports diffing so you can trace history and see impact vs. changes



The screenshot shows the output of the bloaty tool, which is a size profiler for binaries. It displays two columns: FILE SIZE and VM SIZE, each with a percentage and a size in Mi or Ki. The sections are listed on the right side of the table. The total size is 3.90Mi for the file and 2.46Mi for the VM.

FILE SIZE		VM SIZE		
-----	-----	-----	-----	
47.9%	1.87Mi	76.0%	1.87Mi	.text
12.3%	489Ki	0.0%	0	.strtab
11.7%	468Ki	0.0%	0	.debug_info
5.7%	226Ki	9.0%	226Ki	.eh_frame
5.6%	223Ki	0.0%	0	.symtab
4.4%	174Ki	6.9%	174Ki	.rodata
3.3%	130Ki	0.0%	0	.debug_loc
2.4%	95.0Ki	0.0%	0	.debug_line
1.4%	57.8Ki	2.3%	57.8Ki	.data.rel.ro
1.3%	50.8Ki	0.0%	0	.debug_abbrev
0.0%	0	1.8%	44.6Ki	.bss
1.0%	40.1Ki	1.6%	40.1Ki	.gcc_except_table
1.0%	39.3Ki	1.6%	39.3Ki	.eh_frame_hdr
0.6%	25.1Ki	0.0%	0	.debug_ranges
0.6%	24.3Ki	0.0%	0	.debug_str
0.2%	9.20Ki	0.4%	9.20Ki	.data
0.2%	7.53Ki	0.3%	7.53Ki	__libc_freeres_fn
0.2%	7.05Ki	0.0%	0	.note.stapsdt
0.2%	6.36Ki	0.2%	5.56Ki	[25 Others]
0.1%	2.69Ki	0.0%	0	[ELF Section Header]
0.1%	2.53Ki	0.0%	0	.debug_aranges
100.0%	3.90Mi	100.0%	2.46Mi	TOTAL



# Level up your tooling game

(and maybe come over to talk with us – we have some cool stuff for CMake and Bazel [and much more] ;)